

Double-click (or enter) to edit

Objective: To load the heart disease dataset and understand its basic structure. Tasks:

1. Import Pandas and NumPy libraries.
2. Load the dataset using pandas `read_csv()`.
3. Display first five and last five rows.
4. Find number of rows and columns.
5. Display column name

```
import pandas as pd
import numpy as np
```

```
!unzip -o /content/heart_disease.csv.zip -d /content/
```

```
df = pd.read_csv('/content/Heart_Disease_Prediction.csv')
display(df.head())
```

```
unzip:  cannot find or open /content/heart_disease.csv.zip, /content/heart_disea
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	S
0	70	1	4	130	322	0	2	109	0	2.4	
1	67	0	3	115	564	0	2	160	0	1.6	
2	57	1	2	124	261	0	0	141	0	0.3	
3	64	1	4	128	263	0	0	105	1	0.2	
4	74	0	2	120	269	0	2	121	1	0.2	

```
display(df.tail())
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	
265	52	1	3	172		199	1	0	162	0	0.5
266	44	1	2	120		263	0	0	173	0	0.0
267	56	0	2	140		294	0	2	153	0	1.3
268	57	1	4	140		192	0	0	148	0	0.4
269	67	1	4	160		286	0	2	108	1	1.5

```
display(df.head(5))
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	S
0	70	1	4	130		322	0	2	109	0	2.4
1	67	0	3	115		564	0	2	160	0	1.6
2	57	1	2	124		261	0	0	141	0	0.3
3	64	1	4	128		263	0	0	105	1	0.2
4	74	0	2	120		269	0	2	121	1	0.2

```
num_rows, num_cols = df.shape
print(f"Number of rows: {num_rows}")
print(f"Number of columns: {num_cols}")
```

Number of rows: 270
Number of columns: 14

```
print("Column Names:")
for col in df.columns:
    print(col)
```

Column Names:
Age
Sex
Chest pain type
BP
Cholesterol
FBS over 120

EKG results
 Max HR
 Exercise angina
 ST depression
 Slope of ST
 Number of vessels fluro
 Thallium
 Heart Disease

Objective: To identify numerical and categorical variables and their data types. Tasks:

1. Display dataset information using info().
2. Identify numerical and categorical features.
3. Display data types of all columns.
4. Check for missing values.

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              270 non-null    int64  
 1   Sex              270 non-null    int64  
 2   Chest pain type 270 non-null    int64  
 3   BP               270 non-null    int64  
 4   Cholesterol      270 non-null    int64  
 5   FBS over 120     270 non-null    int64  
 6   EKG results      270 non-null    int64  
 7   Max HR           270 non-null    int64  
 8   Exercise angina  270 non-null    int64  
 9   ST depression     270 non-null    float64 
 10  Slope of ST       270 non-null    int64  
 11  Number of vessels fluro 270 non-null  int64  
 12  Thallium          270 non-null    int64  
 13  Heart Disease     270 non-null    object 
dtypes: float64(1), int64(12), object(1)
memory usage: 29.7+ KB
```

```
numerical_features = df.select_dtypes(include=['int64', 'float64'])
categorical_features = df.select_dtypes(include=['object'])

print("Numerical Features:")
print(numerical_features.columns)

print("\nCategorical Features:")
print(categorical_features.columns)

Numerical Features:
Index(['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120',
```

```
'EKG results', 'Max HR', 'Exercise angina', 'ST depression',
'Slope of ST', 'Number of vessels fluro', 'Thallium'],
dtype='object')
```

Categorical Features:

```
Index(['Heart Disease'], dtype='object')
```

```
print("Data Types of All Columns:")
print(df.dtypes)
```

Data Types of All Columns:

Age	int64
Sex	int64
Chest pain type	int64
BP	int64
Cholesterol	int64
FBS over 120	int64
EKG results	int64
Max HR	int64
Exercise angina	int64
ST depression	float64
Slope of ST	int64
Number of vessels fluro	int64
Thallium	int64
Heart Disease	object
dtype: object	

```
print(df.isnull().sum())
```

Age	0
Sex	0
Chest pain type	0
BP	0
Cholesterol	0
FBS over 120	0
EKG results	0
Max HR	0
Exercise angina	0
ST depression	0
Slope of ST	0
Number of vessels fluro	0
Thallium	0
Heart Disease	0
dtype: int64	

Objective: To compute basic descriptive statistics using Pandas and NumPy. Tasks:

1. Display statistical summary using `describe()`.
2. Calculate mean, median, and standard deviation of age.
3. Calculate minimum and maximum cholesterol values.

Double-click (or enter) to edit

```
display(df.describe())
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	
count	270.000000	270.000000	270.000000	270.000000	270.000000	270.000000	27
mean	54.433333	0.677778	3.174074	131.344444	249.659259	0.148148	
std	9.109067	0.468195	0.950090	17.861608	51.686237	0.355906	
min	29.000000	0.000000	1.000000	94.000000	126.000000	0.000000	
25%	48.000000	0.000000	3.000000	120.000000	213.000000	0.000000	
50%	55.000000	1.000000	3.000000	130.000000	245.000000	0.000000	
75%	61.000000	1.000000	4.000000	140.000000	280.000000	0.000000	
max	77.000000	1.000000	4.000000	200.000000	564.000000	1.000000	

Start coding or generate with AI.

```
age_mean = df['Age'].mean()
age_median = df['Age'].median()
age_std = df['Age'].std()

print(f"Mean Age: {age_mean:.2f}")
print(f"Median Age: {age_median:.2f}")
print(f"Standard Deviation of Age: {age_std:.2f}")
```

Mean Age: 54.43
 Median Age: 55.00
 Standard Deviation of Age: 9.11

```
cholesterol_min = df['Cholesterol'].min()
cholesterol_max = df['Cholesterol'].max()

print(f"Minimum Cholesterol: {cholesterol_min}")
print(f"Maximum Cholesterol: {cholesterol_max}")
```

Minimum Cholesterol: 126
 Maximum Cholesterol: 564

▼ Interpretation of Distributions

Age Distribution: The histogram for 'Age' (which you have just plotted) shows a distribution that is somewhat symmetrical, possibly with a slight left-skew. The bulk of the data appears to be concentrated in the middle age ranges, generally between 45 and 65 years. The mean age (54.43) and median age (55.00) are very close, which confirms a relatively symmetrical distribution. This suggests that the heart disease dataset primarily consists of middle-aged to older individuals.

Cholesterol Distribution: The histogram for 'Cholesterol' (which you also plotted) likely indicates a distribution that is moderately right-skewed. This means that while most patients have cholesterol levels within a typical range, there is a tail extending towards higher cholesterol values, indicating a presence of some patients with significantly elevated cholesterol. The mean cholesterol (249.66) being slightly higher than the median (245.00) supports this right-skewness.

Objective: To analyze distributions and frequency of variables. Tasks:

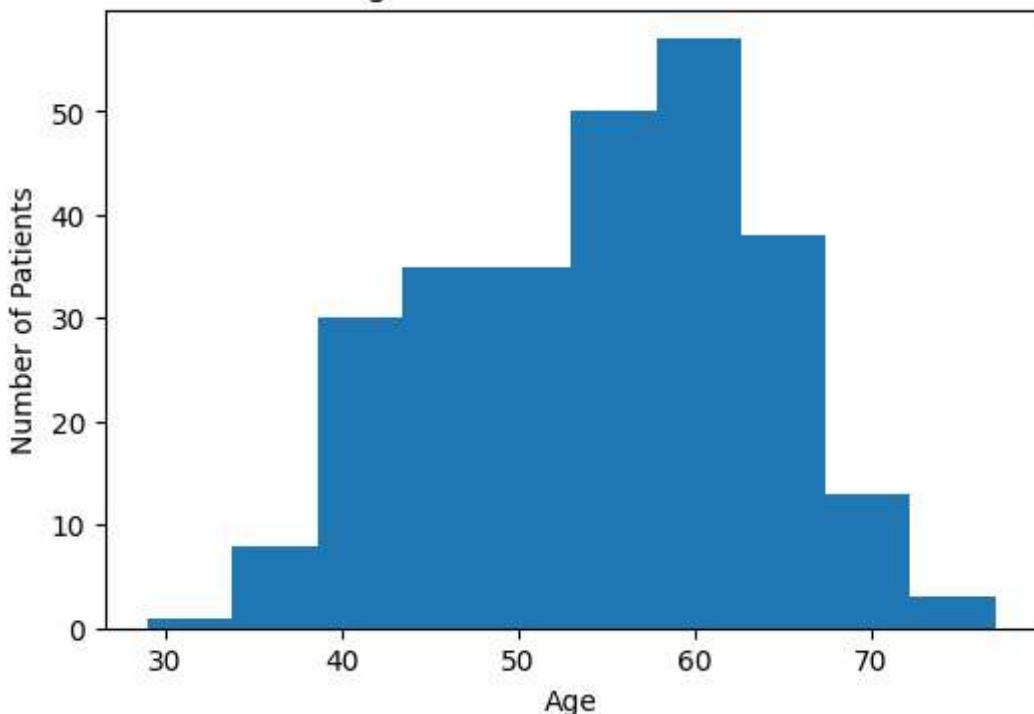
1. Count number of patients with and without heart disease.
2. Plot histogram for age.
3. Plot histogram for cholesterol levels.
4. Interpret distributions.

```
heart_disease_counts = df['Heart Disease'].value_counts()  
print("Number of patients with and without heart disease:")  
print(heart_disease_counts)
```

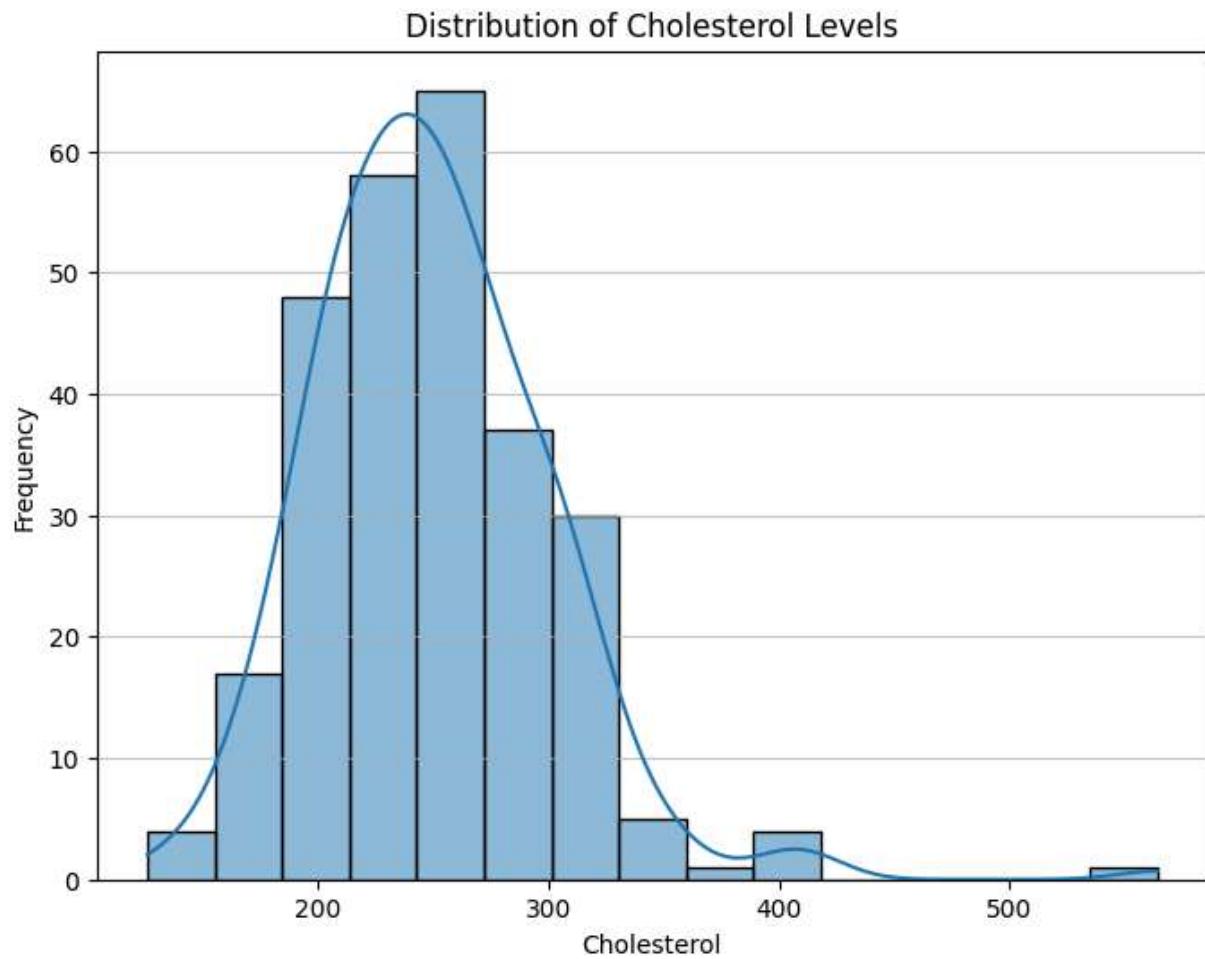
```
Number of patients with and without heart disease:  
Heart Disease  
Absence      150  
Presence     120  
Name: count, dtype: int64
```

```
import matplotlib.pyplot as plt  
import pandas as pd  
  
# Plot histogram for age  
plt.figure(figsize=(6,4))  
plt.hist(df['Age'], bins=10) # Corrected column name to 'Age'  
plt.xlabel("Age")  
plt.ylabel("Number of Patients")  
plt.title("Age Distribution of Patients")  
plt.show()
```

Age Distribution of Patients



```
plt.figure(figsize=(8, 6))
sns.histplot(df['Cholesterol'], kde=True, bins=15)
plt.title('Distribution of Cholesterol Levels')
plt.xlabel('Cholesterol')
plt.ylabel('Frequency')
plt.grid(axis='y', alpha=0.75)
plt.show()
```



Objective: To study relationships between variables using basic exploration. Tasks:

1. Compute correlation matrix.
2. Identify highly correlated variables.
3. Plot scatter plot between age and maximum heart rate.
4. Interpret relationships.

```
correlation_matrix = df.corr(numeric_only=True)  
display(correlation_matrix)
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	res
Age	1.000000	-0.094401	0.096920	0.273053	0.220056	0.123458	0.12
Sex	-0.094401	1.000000	0.034636	-0.062693	-0.201647	0.042140	0.03
Chest pain type	0.096920	0.034636	1.000000	-0.043196	0.090465	-0.098537	0.07
BP	0.273053	-0.062693	-0.043196	1.000000	0.173019	0.155681	0.11
Cholesterol	0.220056	-0.201647	0.090465	0.173019	1.000000	0.025186	0.16
FBS over 120	0.123458	0.042140	-0.098537	0.155681	0.025186	1.000000	0.05
EKG results	0.128171	0.039253	0.074325	0.116157	0.167652	0.053499	1.00
Max HR	-0.402215	-0.076101	-0.317682	-0.039136	-0.018739	0.022494	-0.07
Exercise angina	0.098297	0.180022	0.353160	0.082793	0.078243	-0.004107	0.09
ST depression	0.194234	0.097412	0.167244	0.222800	0.027709	-0.025538	0.12
Slope of ST	0.159774	0.050545	0.136900	0.142472	-0.005755	0.044076	0.16
Number of vessels fluro	0.356081	0.086830	0.225890	0.085697	0.126541	0.123774	0.11
Thallium	0.106100	0.391046	0.262659	0.132045	0.028836	0.049237	0.00

Next steps:

[Generate code with correlation_matrix](#)[New interactive sheet](#)

```
# Set a correlation threshold
correlation_threshold = 0.5

print(f"Highly correlated variable pairs (absolute correlation > {correlation_threshold})")

already_checked = set()
for i in range(len(correlation_matrix.columns)):
    for j in range(i + 1, len(correlation_matrix.columns)):
        col1 = correlation_matrix.columns[i]
        col2 = correlation_matrix.columns[j]
        correlation_value = correlation_matrix.iloc[i, j]
```

Highly correlated variable pairs (absolute correlation > 0.5):
- ST depression and Slope of ST: 0.61

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Max HR', data=df)
plt.title('Scatter Plot of Age vs. Maximum Heart Rate')
plt.xlabel('Age')
plt.ylabel('Maximum Heart Rate')
plt.grid(True)
plt.show()
```

